```
   TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 9001
   inet 172.31.47.158  netmask 255.255.240.0  broadcast 172.31.47.255    ②
   inet6 fe80::4a2:8fff:feea:bbba  prefixlen 64  scopeid 0x20<link>
   ether 06:a2:8f:ea:bb:ba  txqueuelen 1000  (Ethernet)
   RX packets 22  bytes 1641 (1.6 KiB)
   RX errors 0  dropped 0  overruns 0  frame 0
   TX packets 33  bytes 2971 (2.9 KiB)
   TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0
 [...]
```

① The primary network interface uses the private IP address 172.31.33.219.

② The secondary network interface uses the private IP address 172.31.47.158.

Each network interface is connected to a private and a public IP address. You'll need to configure the web server to deliver different websites, depending on the IP address. Your virtual machine doesn't know anything about its public IP address, but you can distinguish the requests based on the private IP addresses.

First you need two websites. Run the following commands on your virtual machine in Sydney via the Session Manager to download two simple placeholder websites:

```
$ sudo -s
$ mkdir /var/www/html/a
$ wget -P /var/www/html/a \
➥ https:/ /raw.githubusercontent.com/AWSinAction/code3/main/chapter03
➥ /a/index.xhtml
$ mkdir /var/www/html/b
$ wget -P /var/www/html/b \
➥ https:/ /raw.githubusercontent.com/AWSinAction/code3/main/chapter03
➥ /b/index.xhtml
```

Next, you need to configure the web server to deliver the websites, depending on which IP address is called. To do so, add a file named a.conf under /etc/httpd/conf.d. The following example uses the editor nano:

```
$ nano /etc/httpd/conf.d/a.conf
```