③ Looks up the AMI with the Imagery worker preinstalled from the region map (see listing 16.10)

④ Disables detailed monitoring of EC2 instances to avoid costs

⑤ Launches virtual machines of type t2.micro to run examples under the Free Tier

⑥ Configures a network interface (ENI) with a public IP address and the security group of the worker

⑦ Each virtual machine will execute this script at the end of the boot process. The script starts the Node.js worker.

⑧ Creates an Auto Scaling group that manages the virtual machines running the Imagery worker

⑨ References the launch template

⑩ The Auto Scaling group will spin up at least one and no more than two EC2 instances.

The Auto Scaling group will replace failed EC2 instances.

Spins up EC2 instances distributed among two subnets: AZs

Adds a Name tag to each instance, which will show up at the Management Console, for example

After all that YAML reading, the CloudFormation stack should be created. Verify the status of your stack like this:

```
$ aws cloudformation describe-stacks --stack-name imagery
{
  "Stacks": [{
    [...]
    "Description": "AWS in Action: chapter 16",
    "Outputs": [{
      "Description": "Load Balancer URL",
      "OutputKey": "EndpointURL",
      "OutputValue":
➥ "http://....us-east-1.elb.amazonaws.com"    ①
    }],
    "StackName": "imagery",
    "StackStatus": "CREATE_COMPLETE"           ②
  }]
}
```

① Copy this output into your web browser.

② Waits until CREATE_COMPLETE is reached