```
    - EFSMountTargetB
    Properties:
      TargetGroupARNs:                                          ③
      - !Ref LoadBalancerTargetGroup
      LaunchTemplate:                                           ④
        LaunchTemplateId: !Ref LaunchTemplate
        Version: !GetAtt 'LaunchTemplate.LatestVersionNumber'
      MinSize: 2                                                ⑤
      MaxSize: 4                                                ⑥
      HealthCheckGracePeriod: 300                               ⑦
      HealthCheckType: ELB                                      ⑧
      VPCZoneIdentifier:                                        ⑨
      - !Ref SubnetA
      - !Ref SubnetB
      Tags:                                                     ⑩
      - PropagateAtLaunch: true
        Value: wordpress
        Key: Name
```

① Creates an Auto Scaling group

② Because the EC2 instances require access to the EFS filesystem, waits until CloudFormation creates the mount targets

③ Registers and unregisters virtual machines on the target group of the load balancer

④ References the latest version of the launch template

⑤ Launches at least two machines and ensures that at least two virtual machines are running, one for each of the two AZs for high availability

⑥ Launches no more than four machines

⑦ Waits five minutes to allow the EC2 instance and web server to start before evaluating the health check

⑧ Uses the ELB health check to monitor the health of the virtual machines

⑨ Launches VMs into two different subnets in two different AZs for high availability

⑩ Adds a tag, including a name for all VMs launched by the ASG

You will learn how to create CloudWatch alarms for scaling in the next example. For now, we are using a target-tracking scaling policy that creates CloudWatch alarms automatically in the background. A target-tracking scaling policy works like the thermostat in your home: you define the target, and the thermostat constantly adjusts the heating power to reach the target. Predefined metric specifications to use with target tracking follow:

- `ASGAverageCPUUtilization`—Scale based on the average CPU usage among all instances within an Auto Scaling group
- `ALBRequestCountPerTarget`—Scale based on the number of requests forwarded from the Application Load Balancer (ALB) to a target
- `ASGAverageNetworkIn` and `ASGAverageNetworkOut`—Scale based on the average number of bytes received or sent