Slowing Down Frontier AI

Why Amodei, Altman, Musk are getting flak

—By **Vyom Ramani**

It has been a little over 4 years since AI burst onto the mainstream and in those 4 years there is one thing I had not seen – agreement. The AI industry as a whole has never managed to agree on LLM performance benchmarks, safety procedures or timelines, or even who actually owns the data they are training the models on. And yet, all it took was one blog post on a weekend in the middle of September for them to be on the same page within hours.



Dario Amodei, CEO of Anthropic, published a blog called "We must Pace the Frontier" on Saturday. A couple of hours later, his fiercest competitors – Sam Altman and Elon Musk – both posted how they agree with him on X. His main argument was that the frontier labs they are running need outside eval and shared safety standards before their models become too smart to control. It is ironic that the thing that made them reach a truce in their race for intelligence was their fear of said intelligence.

The pledge nobody believed

Amodei wants to pace the frontier but how do you exactly achieve that? One option is to stop training their models to be smarter until they are completely secure, which neither GPT 6 Astra nor Claude Mythos was when they were announced. But that's not what Amodei wants, he is clear about it. He wants shared safety standards and external evaluators with employee-level access. Most of all, he wants coordination between the frontier labs. The same labs run by companies who, every other day, poach each other's best engineers. All of this sounds promising until it doesn't.

What I am not buying is that without any enforcement, CEOs of three of the biggest AI labs in the world found religion all on the same weekend. It seems to me and many others, more or less, like a regulatory capture. If you are writing the rules yourself and then handing them to regulators to enforce, you can very easily make sure that the companies that have been dominant so far continue to dominate with something of a safety halo attached. Dirk Auer of Truth on the Market calls out the plan's coordination step as exactly what it looks like – collusion among frontier labs.

David Sacks, Chair of the President's Council of Advisors on Science and Technology, seemingly agreed that there is a need for slowing down but also called Amodei and Altman out. They are the two running the frontier, if their unreleased models are dangerous enough for them to parade that they need to slow down, then they should slow down. They don't require permission. Bernie Sanders, US Senator, is of the opinion that it's not pacing that we need but a pause on advanced AI development and a ban on superintelligence before it is too late. Then you have the President of the USA, Donald Trump, who outright rejects the idea of pacing the frontier saying that they are leading the AI race against China and if they stop, China wins.

What a real pledge looks like

I think that this is not that complicated of an issue to fix. If Amodei, Altman, Musk and Hassabis meant what they have said and really want to slow down the frontier then they should simply sign something. A joint declaration of a kind that is audited by a third party that isn't funded by either of them. Something more concrete than a simple blog post or a three word tweet. A binding commitment that a journalist or the public can hold over their heads after 6 months and ask them if they held their end of the bargain. Nobody's done that. Until then, every "pace the frontier" post is marketing dressed up as a conscience.

This doesn't mean that Amodei is wrong to raise the alarm. There is very much a risk with AI development and the possibility of superintelligence beyond human control. The dystopia was never really going to be a sudden hostile takeover. It starts exactly like this – three billionaires talking about the dangers of building this technology at this pace and then building it at full speed anyway.

tinyurl.com/d1i388d