assassinate you in the morning and take power himself. The hit squad is ready, waiting for his command. Give me the order, though, and I'll liquidate him with a precision strike.'

'But the defense minister is my most loyal supporter,' says the Great Leader. 'Only yesterday he said to me—'

'Great Leader, I know what he said to you. I hear everything. But I also know what he said afterwards to the hit squad. And for months I've been picking up disturbing patterns in the data.'

'Are you sure you were not fooled by deepfakes?'

'I'm afraid the data I relied on is 100 per cent genuine,' says the algorithm. 'I checked it with my special deepfake-detecting sub-algorithm. I can explain exactly how we know it isn't a deepfake, but that would take us a couple of weeks. I didn't want to alert you before I was sure, but the data points converge on an inescapable conclusion: a coup is under way. Unless we act now, the assassins will be here in an hour. But give me the order, and I'll liquidate the traitor.'

By giving so much power to the Surveillance & Security Algorithm, the Great Leader has placed himself in an impossible situation. If he distrusts the algorithm, he may be assassinated by the defence minister, but if he trusts the algorithm and purges the defence minister, he becomes the algorithm's puppet. Whenever anyone tries to make a move against the algorithm, the algorithm knows exactly how to manipulate the Great Leader. Note that the algorithm doesn't need to be a conscious entity to engage in such manoeuvres. As Bostrom's paper-clip thought experiment indicates – and as GPT-4 lying to the TaskRabbit worker demonstrated on a small scale – a nonconscious algorithm may seek to accumulate power and manipulate people even without having any human drives like greed or egotism.

If algorithms ever develop capabilities like those in the thought experiment, dictatorships would be far more vulnerable to algorithmic takeover than democracies. It would be difficult for even a super-Machiavellian AI to seize power in a distributed democratic system like the United States. Even if the AI learns to manipulate the US president, it might face opposition from Congress, the Supreme Court, state governors, the media, major corporations and sundry NGOs. How would the algorithm, for example, deal with a Senate filibuster?