extremist literature', 'befriending known terrorists' and 'having technical knowledge necessary to produce dangerous weapons'. Then they need to *decide* whether a particular individual meets enough of these criteria to be labelled a suspected terrorist. Suppose someone watched a hundred extremist videos on YouTube last month, is friends with a convicted terrorist and is currently pursuing a doctorate in epidemiology in a laboratory containing samples of Ebola virus. Should that person be put on the 'suspected terrorists' list? And what about someone who watched fifty extremist videos last month and is a biology undergraduate?

In Romania in the 1970s only humans could make such decisions. By the 2010s humans were increasingly leaving it to algorithms to decide. Around 2014–15 the US National Security Agency deployed an AI system called Skynet that placed people on a 'suspected terrorists' list based on the electronic patterns of their communications, writings, travel and social media postings. According to one report, that AI system 'engages in mass surveillance of Pakistan's mobile phone network, and then uses a machine learning algorithm on the cellular network metadata of 55 million people to try and rate each person's likelihood of being a terrorist'. A former director of both the CIA and the NSA proclaimed that 'we kill people based on metadata'.[10] Skynet's reliability has been severely criticised, but by the 2020s such technology had become far more sophisticated and has been deployed by a lot more governments. Going over massive amounts of data, algorithms can discover completely new criteria for defining someone as 'suspect' that have previously escaped the notice of human analysts.[11] In the future, algorithms could even create an entire new model for how people are radicalised, just by identifying patterns in the lives of known terrorists. Of course, computers remain fallible, as we shall explore in depth in chapter 8. They may well classify innocent people as terrorists or may create a false model for radicalisation. At an even more fundamental level, it is questionable whether the systems' definition of things like terrorism are objective. There is a long history of regimes using the label 'terrorist' to cover any and all opposition. In the Soviet Union, anyone who opposed the regime was a terrorist. Consequently, when an AI labels someone a 'terrorist' it might reflect ideological biases rather than objective facts. The power to make decisions and invent ideas is inseparable from the capacity to make mistakes. Even if no mistakes are committed, the algorithms' superhuman ability to recognise patterns in an ocean of data can